

SKOPE: A connectionist/symbolic architecture of spoken Korean processing

Geunbae Lee and Jong-Hyeok Lee

Department of Computer Science & Engineering
and Postech Information Research Laboratory
Pohang University of Science & Technology
San 31, Hoja-Dong, Pohang, 790-784, Korea
gblee@vision.postech.ac.kr

Abstract

Spoken language processing requires speech and natural language integration. Moreover, spoken Korean calls for unique processing methodology due to its linguistic characteristics. This paper presents SKOPE, a connectionist/symbolic spoken Korean processing engine, which emphasizes that: 1) connectionist and symbolic techniques must be selectively applied according to their relative strength and weakness, and 2) the linguistic characteristics of Korean must be fully considered for phoneme recognition, speech and language integration, and morphological/syntactic processing. The design and implementation of SKOPE demonstrates how connectionist/symbolic hybrid architectures can be constructed for spoken agglutinative language processing. Also SKOPE presents many novel ideas for speech and language processing. The phoneme recognition, morphological analysis, and syntactic analysis experiments show that SKOPE is a viable approach for the spoken Korean processing.

Introduction

Spoken language processing challenges for integration of speech recognition into natural language processing, and must deal with multi-level knowledge sources from signal level to symbol level. The multi-level knowledge integration and handling increase the technical difficulty of both the speech and the natural language processing. In the speech recognition side, the recognition must be at phoneme-level for large vocabulary continuous speech, and the speech recognition module must provide right level of outputs to the natural language module in the form of not single solution but many alternatives of solution hypotheses. The n-best list (?), word-graph (?), and word-lattice (?) techniques are mostly used in this purpose. The speech recognition module can also ask the linguistic scores from the language processing module in a more tightly coupled bottom-up/top-down hybrid integration scheme (?). In the natural language side, the insertion, deletion, and substitution errors of continuous speech must be compensated by robust parsing and partial parsing tech-

niques, e.g. (?). Often the spoken languages are ungrammatical, fragmentary, and contain non-fluencies and speech repairs, and must be processed incrementally under the time constraints (?).

Most of the speech and natural language systems which were developed for English and other Indo-European languages neglect the morphological processing, and integrate speech and natural language at the word level (?; ?). Often these systems employ a pronunciation dictionary for speech recognition and independent dictionaries for natural language processing. However, for the agglutinative languages such as Korean and Japanese, the morphological processing plays a major role in the language processing since these languages have very complex morphological phenomena and relatively simple syntactic functionality. Unfortunately even the Japanese researchers apply degenerated morphological techniques for the spoken Japanese processing (?; ?). Obviously degenerated morphological processing limits the usable vocabulary size for the system, and word-level dictionary results in exponential explosion in the number of dictionary entries. For the agglutinative languages, we need sub-word level integration which leaves rooms for general morphological processing.

The spoken language processing calls for multi-strategic approaches in order to deal with signal level as well as symbol level information in a symbiotic and unified way. Recent development of connectionist speech recognition (?) and connectionist natural language processing (?) shed lights on the connectionist/symbolic hybrid models of spoken language processing, and some of the researches are already available for English and other Indo-European languages (?; ?). We feel that it is the right time to develop connectionist/symbolic hybrid spoken languages processing systems for the agglutinative languages such as Korean and Japanese.

This paper presents one of the such endeavors, SKOPE (Spoken Korean Processing Engine), that has the following unique features: 1) The connectionist and symbolic techniques are selectively used according to their strength and weakness. The learning capability, fault-tolerant property, and ability of simultaneous inte-

gration of multiple signal-level sources make the connectionist techniques suitable to the phoneme recognition from the speech signals, but the structure manipulation and powerful matching (binding) properties of the symbolic techniques are the better choices for the complex morphological processing of Korean. However, the parallel multiple constraint relaxation capability of the connectionist techniques are applied together with the symbolic structure binding techniques for the syntactic processing. 2) The linguistic characteristics of Korean are fully considered in phoneme recognition, speech and language integration, and morphological/syntactic processing. 3) The SKOPE provides multi-level application program interfaces (APIs) which can utilize the phoneme-level or the morphological level or the syntactic level services for the applications such as spoken language interface, voice information retrieval and spoken language translation.

We hope the experience of SKOPE development provide viable answers to some of the open questions to the speech and language processing, such as 1) how learning and encoding can be synergetically combined in speech and language processing, 2) which aspects of system architecture have to be considered in spoken language processing, especially in connectionist/symbolic hybrid systems, and finally 3) what are the most efficient way of speech and language integration, especially for agglutinative languages.

Characteristics of spoken Korean

This section briefly explains the linguistic characteristics of spoken Korean before describing the SKOPE system. In this paper, Yale romanization is used for representing the Korean phonemes. 1) A Korean word, called *Eojeol*, consists of more than one morphemes with clear-cut morpheme boundaries. 2) Korean is a postpositional language with many kinds of noun-endings, verb-endings, and prefinal verb-endings. These functional morphemes determine the noun's case roles, verb's tenses, modals, and modification relations between Eojeols. 3) Korean is a basically SOV language but has relatively free word order compared to the rigid word-order languages, such as English, except for the constraints that the verb must appear in a sentence-final position. However, in Korean, some word-order constraints do exist such that the auxiliary verbs representing modalities must follow the main verb, and the modifiers must be placed before the word (called head) they modify. 4) The unit of pause in speech (which is called *Eonjeol*) may be different from that of a written text (an Eojeol). The spoken morphological analysis must deal with an Eonjeol since no Eojeol boundary can be provided in the speech. 5) Phonological changes can occur in a morpheme, between morphemes in an Eojeol, and even between Eojeols in an Eonjeol. These changes include consonant and vowel assimilation, dissimilation, insertion, deletion, and contraction. 6) Korean has many rising diphthongs that are very similar to

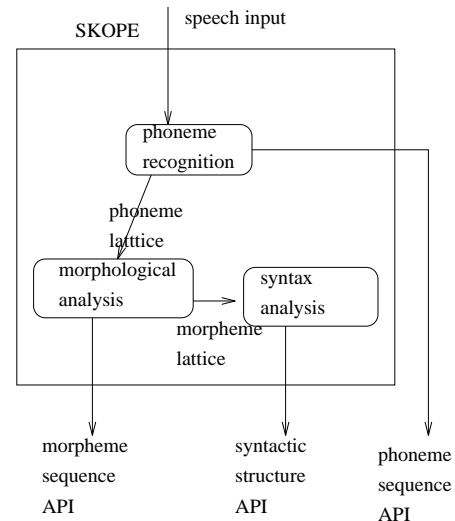


Figure 1: The spoken Korean processing engine architecture. The architecture has two-level interfaces between modules: phoneme lattice and morpheme lattice for efficient and generalized speech and natural language integration.

mono-vowels at signal level. Korean has well-developed syllable structures, and unlike Japanese that has only CV¹ type syllable, Korean has all different types such as CV, VC, V, CVC. Moreover, in CVC type syllable, first and second consonants are almost same in pronunciation. These signal characteristics make it difficult to directly use phonemes or syllables as sub-word recognition units.

The SKOPE architecture

The above spoken Korean characteristics and the relative strength and weakness of symbolic/connectionist techniques result in the general SKOPE architecture which is shown in figure 1. The architecture consists of three different but closely interrelated modules: phoneme recognition, morphological analysis, and syntactic analysis module. The phoneme recognition module processes the signal-level information, and changes it to the symbol-level information (phoneme lattice). The morphological analysis begins the primitive language processing, and connects the speech recognition to the language processing at the phoneme-level. The syntactic analysis module finishes the language processing², and produces the domain independent syntactic structures for application systems. The following subsections briefly describe each module.

¹C: consonant, V: vowel

²We believe that the semantic and pragmatic processing should be integrated into the domain knowledge for *practical application under the current NLP technology*, so we excluded the semantic and pragmatic processing from our general model.

diphone types	diphone numbers	diphone examples
V	21	a, o, wu, i, u, ye,
C1V	378	ha, sa, ka, la, ma, kha,
VC2	147	an, am, eng, em, wun, in,
C2C1	126	ngs, nn, ngt, ngh,

Figure 2: Four different Korean diphone types (V: vowel, C1: syllable-first consonant, C2: syllable-final consonant)

Diphone-based connectionist phoneme recognition

The phoneme recognition is performed by developing the hierarchically organized group of TDNNs (time delay neural networks) (?). Considering the signal characteristics of the Korean phonemes, we define diphones as a new sub-word recognition unit. The defined diphones are shown in figure 2, and are classified into four different types. The diphones have the co-articulation handling features similar to the popular triphones (?) but are much fewer in numbers.

Figure 3 shows the architecture of the component TDNNs in the phoneme recognition module. The whole module consists of total 19 different TDNNs for recognition of the defined Korean diphones. The top-level TDNN identifies the 18 vowel groups of diphones (we reclassified the total 672 diphones into 18 different groups according to the vowels that are contained in the diphones). The 18 different sub-TDNNs recognize the target diphones.

For the training of TDNNs, we manually segment the digitized speech into 200 msec range (which includes roughly left-context phoneme, target diphone, and right context phoneme), and perform 512 order FFTs and 16 step mel-scaling (?) to get the filter-bank coefficients. Each frame size is 10 msec, so 20 (frames) by 16 (mel-scaling factor) values are fed to the TDNNs with the proper output symbols, that is, vowel group name or target diphone names. After the training of each TDNN, the phoneme recognition is performed by feeding 200 msec signals to the vowel group identification network and subsequently to the proper diphone recognition network. The 200 msec signals are shifted by 30 msec steps and continuously fed to the networks to process the continuous speech in an Eonjeol. From the resulting diphone sequences, the necessary phoneme lattice has to be constructed. We use a simple deterministic decoding heuristics and try to maintain all the possible diphone spotting results since the later phonological/morphological processing can safely prune the incorrect recognitions. The decoding begins by grouping the diphones into the same types (see figure 2). The frequency count for each diphone, that is, the number of specific diphones per 30 msec frame shift, is utilized




Figure 3: (a) The TDNN architecture for the vowel group identification. Note the cc group contains no vowels. (b) The architecture of sub-TDNNs for /a/ vowel group. The other 17 sub-TDNNs have the same architecture, but different number of output units according to the number of diphones in each of the vowel group.

to fix the insertion errors by deleting the lower frequency count diphones, and finally the diphones are split into the constituent phonemes by merging the same phonemes in the neighboring diphones.

Table-driven morphological and phonological analysis

The morphological analysis starts with the phoneme lattice. The phoneme lattice delivers the alternative phonetic transcriptions³ of input speech, which must be searched by the morphological/phonological analyzer to reconstruct the orthographic morpheme strings. The conventional morphological analysis procedure (?), that is, morpheme segmentation, morphotactics modeling, and orthographic rule (or phonological rule) modeling, must be augmented and extended as the followings: 1) The conventional morpheme segmentation is extended to deal with the exponential number of phoneme sequences and between-morpheme phonological changes during the segmentation, 2) the morphotactics modeling is extended to cope with the complex verb and noun-endings (or postpositions), and 3) the orthographic rule modeling is combined with the phonological rule modeling to correctly transform the phonetic transcriptions to the orthographic morpheme sequences.

The central part of the morphological analysis lies in the dictionary construction. In our dictionary, each

³Unlike English, the Korean alphabet is truly phonetic in the sense that each phoneme is pronounced as it is written. That is why we sometimes use *phonetic* and *phonemic* interchangeably.

phonetic transcription header	original morpheme	left morphological connectivity	right morphological connectivity	left phonemic connectivity	right phonemic connectivity
ci-wu	ci-wu	regular verb	regular verb	'c' sound no-change	'wu' sound no-change
l	l	adnominalizing verb-ending	adnominalizing verb-ending	'l' sound no-change	'l' sound no-change
sswu	swu	bound-noun	bound-noun	's' sound change to 'ss'	'wu' sound no-change

Figure 4: The morpheme-level phonetic dictionary.




Figure 5: Morphological parsing of the phoneme lattice (from top: output morpheme sequence in an Eonjeol, triangular parsing table, input phoneme lattice).

phonetic transcription of single morpheme has a separate dictionary entry. Figure 4 shows the unified dictionary both for speech and language processing (called morpheme-level phonetic dictionary) with three different morpheme entries *ci-wu*, *l*, *swu*.

The extended morphological analysis is based on the well-known tabular parsing technique for context-free language (?) and augmented to handle the Korean phonological rules and phoneme-lattice input. Figure 5 shows our extended table-driven morphological analysis process. The example phoneme lattice was obtained from the input speech *ci-wul-sswu* (removable), and the morphological analysis produces *ci-wu+l+swu* (remove+ADNOMINAL+BOUND-NOUN), where '+' is the morpheme boundary, and '-' is the syllable boundary.

The extended morpheme segmentation is basically performed using the dictionary search. During the left-

to-right scan of the input phoneme lattice, when a morpheme boundary is found in the lattice, the morpheme is enrolled in the triangular table in an appropriate position. For example, in figure 5, morphemes such as *ci-wu*, *l*, *swu*, etc are enrolled in the table position (1,3), (4,4), (5,6), etc. The position (i,j) designates the starting and ending position of the enrolled morphemes. However since the input is a phoneme-lattice, total exponential time is required to find all the possible morpheme boundaries. To cope with such exponential explosion, the dictionary is organized as trie structure (?) using the phonetic transcriptions as trie indices, and breadth-first search of the trie can prune the unnecessary phoneme sequences earlier in the search.

The morphotactics modeling is necessary after all the morphemes are enrolled in the table in order to combine only legal morphemes into an Eojeol (Korean word), and the process is called morpheme-connectivity-checking. Since Korean has well developed postpositions (noun-ending, verb-ending, prefinal verb-ending) which play as grammatical functional morphemes, we must assign each morpheme proper part-of-speech (POS) tags for the efficient connectivity checking. Our more than 200 POS tags which are refined from the 13 major Korean lexical categories are hierarchically organized, and contained in the dictionary (in the name of morphological connectivity, see figure 4). In the case of idiomatic expressions, we place such idioms directly in the dictionary for efficiency, where two different POS tags are necessary for the left and the right morphological connectivity. For single morpheme, the left and the right POS tags are always same. The separate morpheme-connectivity-matrix indicates the legal morpheme combinations, and the morphotactics modeling is performed using the POS tags (in the dictionary) and morpheme-connectivity-matrix.

The orthographic rule modeling must be integrated with the phonological rule modeling in spoken language processing. Since we must deal with the phoneme lattice, the conventional rule-based modeling requires exponential number of rule application (?). So our solution is based on the declarative modeling of both orthographic and phonological rules in uniform way. That is, in our dictionary, the conjugated verb forms as well as the original verb forms are enrolled, and the same morphological connectivity information is applied. The phonological rule modeling is also accomplished declaratively by having the phonemic connectivity information in the dictionary. The phonemic connectivity information for each morpheme declares the possible phonemic changes in the first (left) and last (right) positioned phonemes in the morpheme, and the separate phoneme-connectivity-matrix indicates the legal sound combinations in Korean phonology. For example, in figure 5, the morpheme *l* can be combined with the morpheme *swu* during the morpheme connectivity checking even if *swu* is actually pronounced as *sswu* because the phoneme-connectivity-matrix supports the legality of the combi-

nation of *l* sound with *ss* sound⁴. In this way, we can declaratively model all the major Korean phonology rules such as second consonant standardization, consonant assimilation, palatalization, glotalization, insertion, deletion, and contraction.

Table-driven connectionist/symbolic syntax analysis

The phoneme lattice-based morphological analysis produces the morphologically analyzed (segmented and stem reconstructed) morpheme sequences. Since there are usually more than one analysis results due to the errors of speech recognition process, the outputs are usually organized as morpheme lattice. For the seamless integration of the morphological analysis with the syntax analysis, we employ the same table-driven control for the syntax analysis as well as the morphological analysis.

We extend the category formation and functional application rules in the previous categorial unification grammar($?$; $?$) to deal with the word order variations in Korean:

- if category $a \in C$, then $a \in C'$
- if category $a \in C'$, and category set $S \in C'$, then $a/S \in C'$ and $a \setminus S \in C'$

where S is an unordered set of categories.

- left cancellation: $a_i \setminus \{a_1, a_2, \dots, a_n\}$ results in $b \setminus \{a_1, a_2, \dots, a_{i-1}, a_{i+1}, \dots, a_n\}$
- right cancellation: $b / \{a_1, a_2, \dots, a_n\} a_i$ results in $b / \{a_1, a_2, \dots, a_{i-1}, a_{i+1}, \dots, a_n\}$

The syntax analysis is performed by interactive relaxation (spreading activation) parsing on the categorial grammar where the position of the functional applications are controlled by a triangular table. The original interactive relaxation parsing ($?$) was extended to provide efficient constituent searching and expectation generation through positional information provided by categorial grammar and triangular table. Figure 6 shows table-driven interactive relaxation parsing.

The interactive relaxation process consists of the following three steps that are repetitively executed: 1) add nodes, 2) spread activation, and 3) decay.

add nodes Grammar nodes (syntactic categories from the dictionary) are added for each sense of the morphemes when the parsing begins. A grammar node which has more activation than the predefined threshold Θ generates new nodes in the proper positions (to be discussed shortly). The newly generated nodes search for the constituents (expectations) which are in the appropriate table positions, and are of proper function applicable categories. For example, in figure 6, when $np \setminus np(2,2)$ fires, it generates

⁴This legality comes from the Korean phonology rule *glotalization* (one form of consonant dissimilation) stating that *s* sound becomes *ss* sound after *l* sound.




Figure 6: Table-driven interactive relaxation parsing of a categorial grammar. The input sentence is *phai-l-tul-ul ciwu-ela* (delete the files). Only single morpheme chain and only one sense for each morpheme is shown as input for clear illustration. The *subj*, *obj*, *comm* indicate $np[subj]$, $np[obj]$, $s[command]$, respectively. The table contains only the nodes that participate in the final parse trees.

$np(1,2)$. The generated $np(1,2)$ searches for the constituents $np(1,1)$ to be combined with $np \setminus np(2,2)$.

spread activation The bottom-up spreading activation is as follows:

$$n \times \rho \times a \times \frac{a_i^2}{\sum_{j=1}^n a_j^2}$$

where predefined portion ρ of total activation a is passed upward to the node with activation a_i among the n parents each with node activation a_j . In other words, the node with large activation gets more and more activation, and it gives an inhibition effects without explicit inhibitory links ($?$). The top-down spreading activation uniformly distributes:

$$\rho' \times a$$

among the children where ρ' is predefined portion of the source activation a .

decay The node's activation is decayed with time. The node with less constituents than needed gets penalties plus decays:

$$a \times (1 - d) \times \frac{C_a}{C_r}$$

where a is an activation value, d is a decay ratio, and C_a , C_r is the actual and required constituents. After the decay, the node with less activation than the predefined threshold Φ is removed from the table.

The node generation and constituent search positions are controlled by the triangular table. When the node $a(i,j)$ acts as an argument, it generates node only in the position (k,j) where $1 < k < j$, and the generated node searches for the constituents (functors) only in the position $(k,i-1)$. Or when the node is generated in the

position (i,k) where $j < k < \text{number-of-morphemes}$, it searches for the position (j+1,k) for its constituents. When the node acts as a functor, the same position restrictions also apply for the node generation and the argument searching. The position control combined with the interactive relaxation guarantees an efficient, lexically oriented, and robust syntax analysis of spoken languages.

Implementation and experiments

The SKOPE was fully implemented in UNIX/C platform, and have been extensively tested in practical domains such as natural language interface to operating systems. The phoneme recognition module targets 1000 morpheme continuous speech, currently speaker dependent due to the short of standard speech database for Korean. The unified morpheme-level phonetic dictionary has about 1000 morpheme entries and compiled into the trie structure. The morpheme-connectivity-matrix and phoneme-connectivity-matrix are encoded with the special Korean POS (part-of-speech) symbols and compressed.

This section demonstrates the SKOPE's performance in continuous diphone recognition, morphological analysis, and syntax analysis experiments. For the continuous diphone recognition experiment, we generated about 5500 diphone patterns from the 990 Eojeol patterns (66 Eojeols, 15 times pronunciation) for the training of TDNNs. In the performance phase, the new 2600 test Eojeol patterns (260 Eojeol, 10 times pronunciation) are continuously shifted with 30 msec step, and generate 7772 test diphone patterns disjoint from the training patterns. Figure 7-a shows the continuous diphone recognition performance. The *correct* designates that the correct target diphones were spotted in the testing position, and the *delete* designates the other case. The *insert* designates that the non-target diphones were spotted in the testing position. To compare the ability of handling the continuous speech, we also tested the diphone recognition using the hand-segmented test patterns with the same 7772 target diphones. Figure 7-b shows the segmented diphone recognition performance. Since the test data are already hand-segmented before input, there are no insertion and deletion errors in this case. The fact that the segmented speech performance is not much better than the continuous one (93.8% vs. 93.4%) demonstrates the diphone's suitability to handling the continuous speech.

For the morphological analysis performance, we used the same 990 Eojeol patterns to train the phoneme recognition module, and the 2600 Eojeol patterns to test the morphological analysis performance directly from the speech input. Figure 8 shows the results.

This experiment shows that most of the morphological errors are propagated from the incorrect (deleted) or spurious (inserted) phoneme recognition results. To see the original performance of the morphological and syntactic analysis modules assuming no speech recognition

(a) continuous diphones				
	total	correct	delete	insert
pattern size (rec. rate)	7772	7259 (93.4%)	513 (6.6%)	3000 (38.6%)

(b) segmented diphones			
	vowel group	sub-TDNNs average	total average
rec. rate	94.8%	98.2%	93.8%

Figure 7: (a) Continuous diphone recognition versus (b) segmented diphone recognition

	total	correct	deleted	inserted
number of morphemes	9605	7696 (80.1%)	1909 (19.8%)	7182

Figure 8: Morphological analysis from continuous speech signals. The table indicates that, among the total 9605 morphemes in 2600 Eojeol patterns, the 80.1% are correctly recognized and analyzed, and 19.8% cannot be analyzed for deletion errors. The 7182 spurious morphemes are also generated due to the speech insertion errors.

error, we artificially made the phoneme lattices by mutating the correctly recognized phoneme sequences according to the phoneme recognizer's confusion matrix. Each phoneme lattice was made to contain at least one correct recognition result, so the phoneme recognition performance is assumed to be perfect except the artificially made insertion errors (mutations). In this way, we made 6 or 7 lattices for each of the 50 sentences, altogether 330 phoneme lattices. The average phoneme alternatives per single correct phoneme in the lattice are 2.3, and average sentence length is 31 phonemes. This means there are average 2.3^{31} phoneme chains in each lattice. The used sentences are natural language commands to UNIX (?) and are fairly complex which have one or two embedded sentences or conjunctions. Figure 9 shows the morphological and syntactic analysis results for these artificially made phoneme lattices. For the syntactic level interactive relaxation, we used the following parameters (which are experimentally determined): upward propagation portion ρ 0.05, downward propagation portion ρ' 0.03, decay ratio d 0.87, the node generation threshold Θ 0.51, and the node removal threshold Φ 0.066.

The morphological analysis was perfect as shown in the table. Since the phoneme lattice was made to contain at least one correct phoneme recognition result, the morphological analysis must be perfect as long as the morpheme is enrolled in the dictionary and the connectivity information can cover all the morpheme combinations. This was possible due to the small number of tested sentences (50 sentences). This results verify

	total sentences	correctly analyzed sentences
morphological	330	330
syntactic (1-best)	330	117 (35.5%)

Figure 9: The morphological and syntactic analysis from the artificially made phoneme lattices.

that most of the morphological analysis errors from real speech input are actually propagated from the phoneme recognition errors as discussed before. However, the syntax analysis results are marginal here since we only count the single best scored tree, and we don't use yet any semantic feature in the analysis. The syntax analysis failures mainly come from 1) the insertion errors (artificial mutations) in the phoneme lattices⁵, which result in ambiguous morpheme lattice, and finally produce redundant syntax trees, and 2) the inherent structural ambiguities in the sentence. These failures should be greatly reduced if we generate n-best scored parse trees, and let the semantic processing module select the correct ones as is usually done in most of the probabilistic parsing schemes(?).

Conclusions and future works

This paper explains the design and implementation of spoken Korean processing engine, which is a connectionist/symbolic hybrid model of spoken language processing by utilizing the linguistic characteristics of Korean. The SKOPE model demonstrates the synergistic integration of connectionist and symbolic techniques by considering the relative strength and weakness of two different techniques, and also demonstrates the phoneme level speech and language integration for general morphological processing for agglutinative languages. Besides the above two major contributions, the SKOPE architecture has the following unique features in spoken language processing: 1) the diphones are newly developed as a sub-word recognition unit for connectionist Korean speech recognition, 2) the morphological and syntactic analysis are tightly coupled by using the uniform table-driven control, 3) the phonological and orthographic rules are uniformly co-modeled declaratively, and 4) the table-driven interactive relaxation parsing and extension of the categorial grammar can provide robust handling of word-order variations in Korean.

However, current implementation of the system still suffers from excessive continuous speech recognition errors. Since the large vocabulary continuous speech recognition is still an open problem, we cannot hope for the 100% correct speech recognition results in the near future. Currently, we are pursuing multi-strategic

approaches to the advanced spoken language processing model, including optimizing TDNN-based phoneme recognition module, integrating HMM-based morpheme recognition module into the connectionist phoneme recognition, and incorporating probabilistic searches into the morphological analysis process as well as the syntactic analysis process. We are also developing applications on top of our SKOPE, including speech-to-speech translation system and intelligent interface agent for UNIX operating system. We hope our approach could be extended to other agglutinative languages such as Japanese, Finnish, and Turkish, and also to the languages that have complex morphological phenomena such as German and Dutch.

Acknowledgments

This research was supported partly by a grant from KOSEF (Korea Science and Engineering Foundation) and PIRL (Postech Information Research Laboratory). The SKOPE's various modules were programmed by our students: the phoneme recognition module by Kyunghee Kim & Kyubong Baac, the morphological analysis module by EunChul Lee & Wonil Lee, and finally the syntax analysis module by Wonil Lee.

⁵Recall we generated average 2.3 phonemes per single correct phoneme.